




# ADV DATA MINING AND ANALYTICS


## Regularization & GLM

Rouzbeh Razavi, PhD


## Our Old Friend !




# Overfitting: Regression



$\rightarrow \theta_0 + \theta_1 x$   
"Underfit" "High bias"



$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2$




$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$   
"Overfit" "High variance"

To recap, if we have too many features then the learned hypothesis may give a cost function of exactly zero

- But this tries too hard to fit the training set
- Fails to provide a *general* solution - **unable to generalize** (apply to new examples)


# Overfitting: Classification




$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

( $g$  = sigmoid function)

**UNDERFITTING**  
(high bias)



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2)$$



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \theta_6 x_1^3 x_2 + \dots)$$

**OVERFITTING**  
(high variance)

# Overfitting: Counter Measures

- How do we deal with this?
  - 1) **Reduce number of features**
    - Manually select which features to keep
    - Model selection algorithms are discussed later (good for reducing number of features)
    - But, in reducing the number of features we lose some information
      - Ideally select those features which minimize data loss, but even so, some info is lost
  - 2) **Regularization**
    - Keep all features, but reduce magnitude of parameters  $\theta$
    - Works well when we have a lot of features, each of which contributes a bit to predicting  $y$



## Cost function optimization for regularization

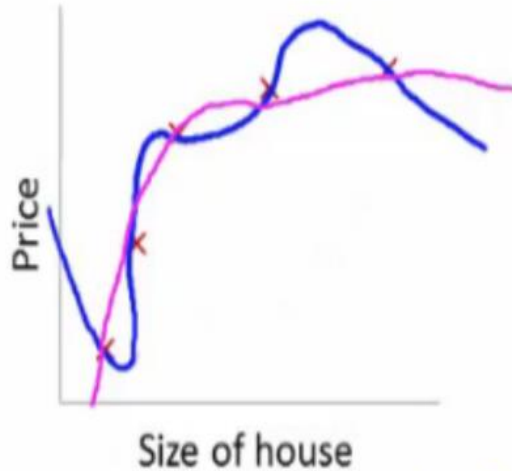
- Penalize and make some of the  $\theta$  parameters really small
  - e.g. here  $\theta_3$  and  $\theta_4$

$$\min_{\theta} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + 1000 \theta_3^2 + 1000 \theta_4^2$$

- The addition in blue is a modification of our cost function to help penalize  $\theta_3$  and  $\theta_4$

# Overfitting: Counter Measures

- So here we end up with  $\theta_3$  and  $\theta_4$  being close to zero (because the constants are massive)
- So we're basically left with a quadratic function



$$\theta_0 + \theta_1 x + \theta_2 x^2 + \cancel{\theta_3 x^3} + \cancel{\theta_4 x^4}$$

The equation shows a polynomial fit. The terms  $\theta_3 x^3$  and  $\theta_4 x^4$  are crossed out with blue X's, and pink arrows point to them from below, indicating they are being discarded or set to zero to counter overfitting.

# Regularization

- In this example, we penalized two of the parameter values
  - More generally, regularization is as follows
- Regularization
  - Small values for parameters corresponds to a simpler hypothesis (you effectively get rid of some of the terms)
  - A simpler hypothesis is less prone to overfitting
- Another example
  - Have 100 features  $x_1, x_2, \dots, x_{100}$
  - Unlike the polynomial example, we don't know what are the high order terms
    - How do we pick the ones to pick to shrink?
  - With regularization, take cost function and modify it to shrink all the parameters
    - Add a term at the end
      - This regularization term shrinks every parameter
      - By convention you don't penalize  $\theta_0$  - minimization is from  $\theta_1$  onwards



# Regularization

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

$\theta_1, \theta_2, \theta_3, \dots, \theta_{100}$

- In practice, if you include  $\theta_0$  has little impact
- $\lambda$  is the **regularization parameter**
  - Controls a trade off between our two goals
    - 1) Want to fit the training set well
    - 2) Want to keep parameters small
- With our example, using the **regularized objective** (i.e. the cost function with the regularization term) you get a much smoother curve which fits the data and gives a much better hypothesis

# Regularization

- If  $\lambda$  is very large we end up penalizing ALL the parameters ( $\theta_1$ ,  $\theta_2$  etc.) so all the parameters end up being close to zero
  - If this happens, it's like we got rid of all the terms in the hypothesis
    - This results here is then underfitting
  - So this hypothesis is too biased because of the absence of any parameters (effectively)

So,  $\lambda$  should be chosen carefully - not too big...

- We look at some automatic ways to select  $\lambda$  later in the course

# Regularization


- If  $\lambda$  is very large we end up penalizing ALL the parameters ( $\theta_1$ ,  $\theta_2$  etc.) so all the parameters end up being close to zero
  - If this happens, it's like we got rid of all the terms in the hypothesis
    - This results here is then underfitting
  - So this hypothesis is too biased because of the absence of any parameters (effectively)

So,  $\lambda$  should be chosen carefully - not too big...

- We look at some automatic ways to select  $\lambda$  later in the course


# Example: Logistic Regression

Training  
Data



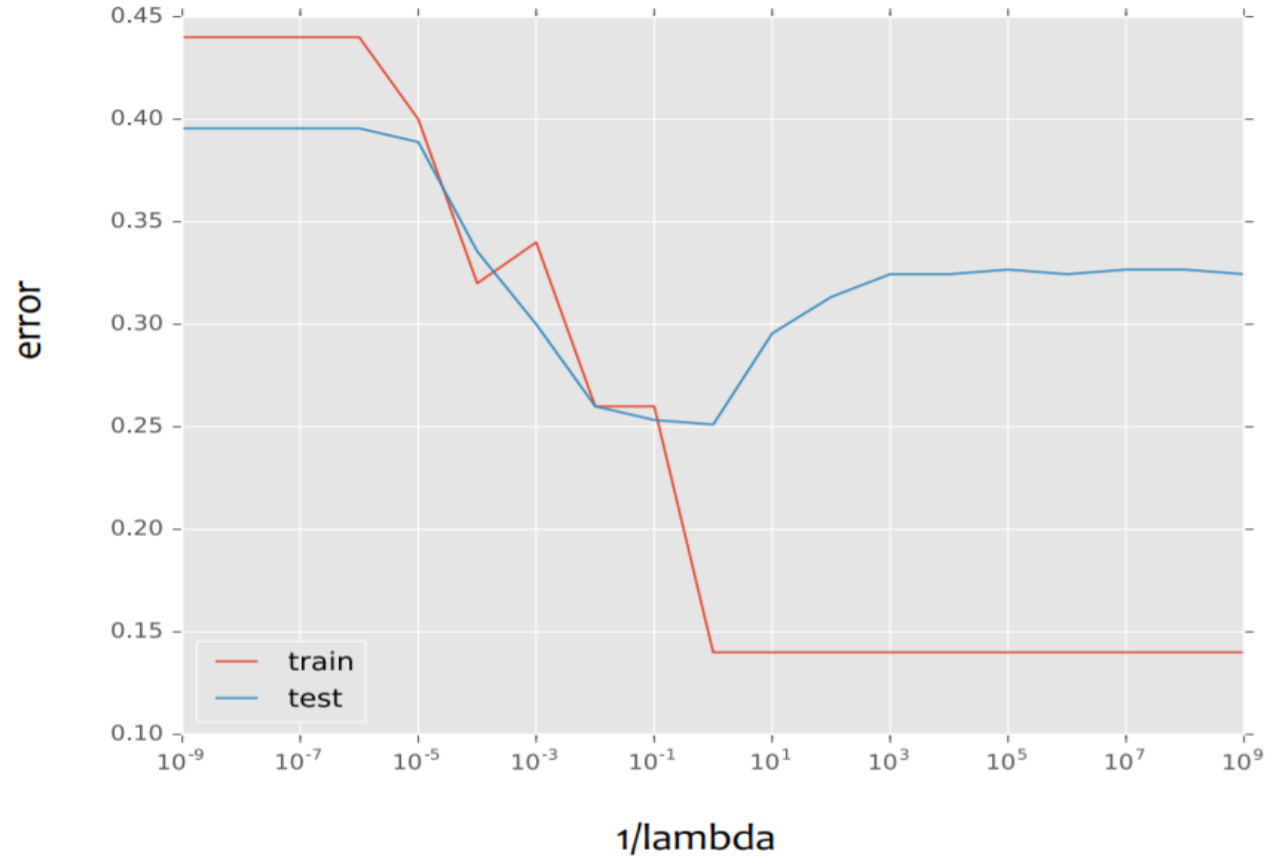
# Example: Logistic Regression

Test  
Data







# Example: Logistic Regression




# Example: Logistic Regression



# Example: Logistic Regression




# Example: Logistic Regression




# Example: Logistic Regression

Classification with Logistic Regression ( $\lambda=1$ )

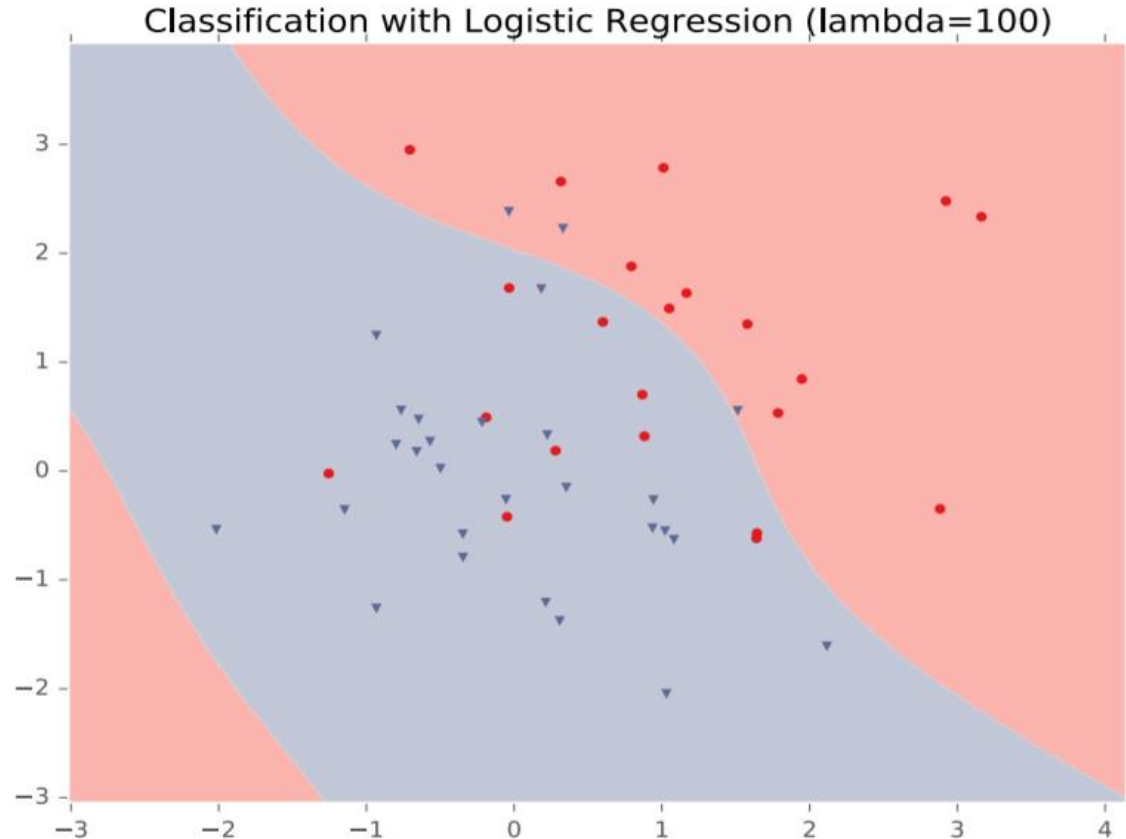





# Example: Logistic Regression



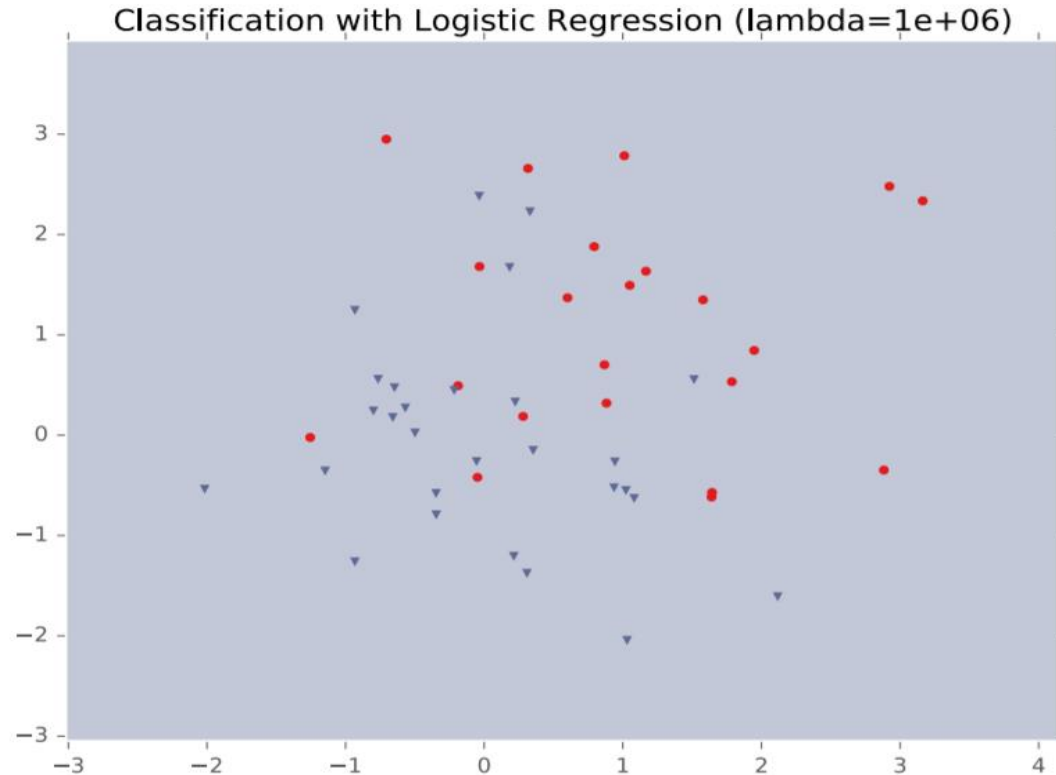
# Example: Logistic Regression



# Example: Logistic Regression




# Example: Logistic Regression




## Another Example

$\lambda = 0$




overfitting


$\lambda = 0.0001$



$\lambda = 0.01$



$\lambda = 1$



underfitting





## Loss Functions to Minimize

Square of Errors  
as seen before

Ridge Regression

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

$\theta_1, \theta_2, \theta_3, \dots, \theta_{100}$

Lasso Regression

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n |\theta_j| \right]$$

$\theta_1, \theta_2, \theta_3, \dots, \theta_{100}$

# Lasso vs Ridge

Note the name “lasso” is actually an acronym for: Least Absolute Selection and Shrinkage Operator

The only difference between the lasso problem and ridge regression is that the latter uses a (squared)  $\ell_2$  penalty  $\|\theta\|_2^2$ , while the former uses an  $\ell_1$  penalty  $\|\theta\|_1$ . But even though these problems look similar, their solutions behave very differently


# Lasso regression

For  $\lambda$  in between these two extremes, we are balancing two ideas: fitting a linear model of  $y$  on  $X$ , and shrinking the coefficients. But the nature of the  $\ell_1$  penalty causes some coefficients to be shrunk to **zero exactly**


This is what makes the lasso substantially different from ridge regression: it is able to perform **variable selection** in the linear model. As  $\lambda$  increases, more coefficients are set to zero (less variables are selected), and among the nonzero coefficients, more shrinkage is employed



# Ridge regression



# Lasso regression





# Lasso regression

As with ridge regression, the penalty term  $\|\boldsymbol{\beta}\|_1 = \sum_{j=1}^p |\beta_j|$  is not fair if the predictor variables are **not on the same scale**. Hence, if we know that the variables are not on the same scale to begin with, we **scale** the columns of  $X$  (to have sample variance 1), and then we solve the lasso problem

## Constrained form

It can be helpful to think of our two problems **constrained form**:



$$\hat{\Theta}^{\text{ridge}} = \underset{\Theta \in \mathbb{R}^p}{\operatorname{argmin}} \|y - X\Theta\|_2^2 \quad \text{subject to } \|\Theta\|_2^2 \leq t$$

$$\hat{\Theta}^{\text{lasso}} = \underset{\Theta \in \mathbb{R}^p}{\operatorname{argmin}} \|y - X\Theta\|_2^2 \quad \text{subject to } \|\Theta\|_1 \leq t$$

Now  $t$  is the tuning parameter (before it was  $\lambda$ ). For any  $\lambda$  and corresponding solution in the previous formulation (sometimes called **penalized form**), there is a value of  $t$  such that the above constrained form has this same solution

In comparison, the usual linear regression estimate solves the **unconstrained** least squares problem; these estimates constrain the coefficient vector to lie in some geometric shape centered around the origin. This generally reduces the variance because it keeps the estimate close to zero. But which shape we choose **really matters!**

# Why does the lasso give zero coefficients?



# Time for Implementation ! Glmnet Package

[https://web.stanford.edu/~hastie/glmnet/glmnet\\_alpha.html](https://web.stanford.edu/~hastie/glmnet/glmnet_alpha.html)

**Rob Tibshirani**



**Trevor Hastie**

ANY  
QUESTIONS  
?